

```
    execve(args[0], args, environ);  
}
```

```
*CTF{simple_simpler_simplest_stack_overflow}
```

misc/Alice's warm up

We are provided with a `data.pk1`, 6 serialized tensors in the `data` folder and a `hint.py` with the following contents:

```
1 import string  
2 assert len(flag)==16  
3 assert flagset=string.printable[0:36]+"*CTF{ALIZE}"
```


Looking at the `data.pk1`, we can deduce that the model is equivalent to three linear layers:

```
1 v0 = torch.tensor(np.frombuffer(open('./data/0', 'rb').read(), dtype=np.float32).re  
2 v1 = torch.tensor(np.frombuffer(open('./data/1', 'rb').read(), dtype=np.float32).re  
3 v2 = torch.tensor(np.frombuffer(open('./data/2', 'rb').read(), dtype=np.float32).re  
4 v3 = torch.tensor(np.frombuffer(open('./data/3', 'rb').read(), dtype=np.float32).re  
5 v4 = torch.tensor(np.frombuffer(open('./data/4', 'rb').read(), dtype=np.float32).re  
6 v5 = torch.tensor(np.frombuffer(open('./data/5', 'rb').read(), dtype=np.float32).re  
7  
8 def forward(x):  
9     z0 = (v0 @ x) + v1  
10     z1 = (v2 @ z0) + v3  
11     z2 = (v4 @ z1) + v5  
12     x = torch.sigmoid(z2)  
13     return x
```

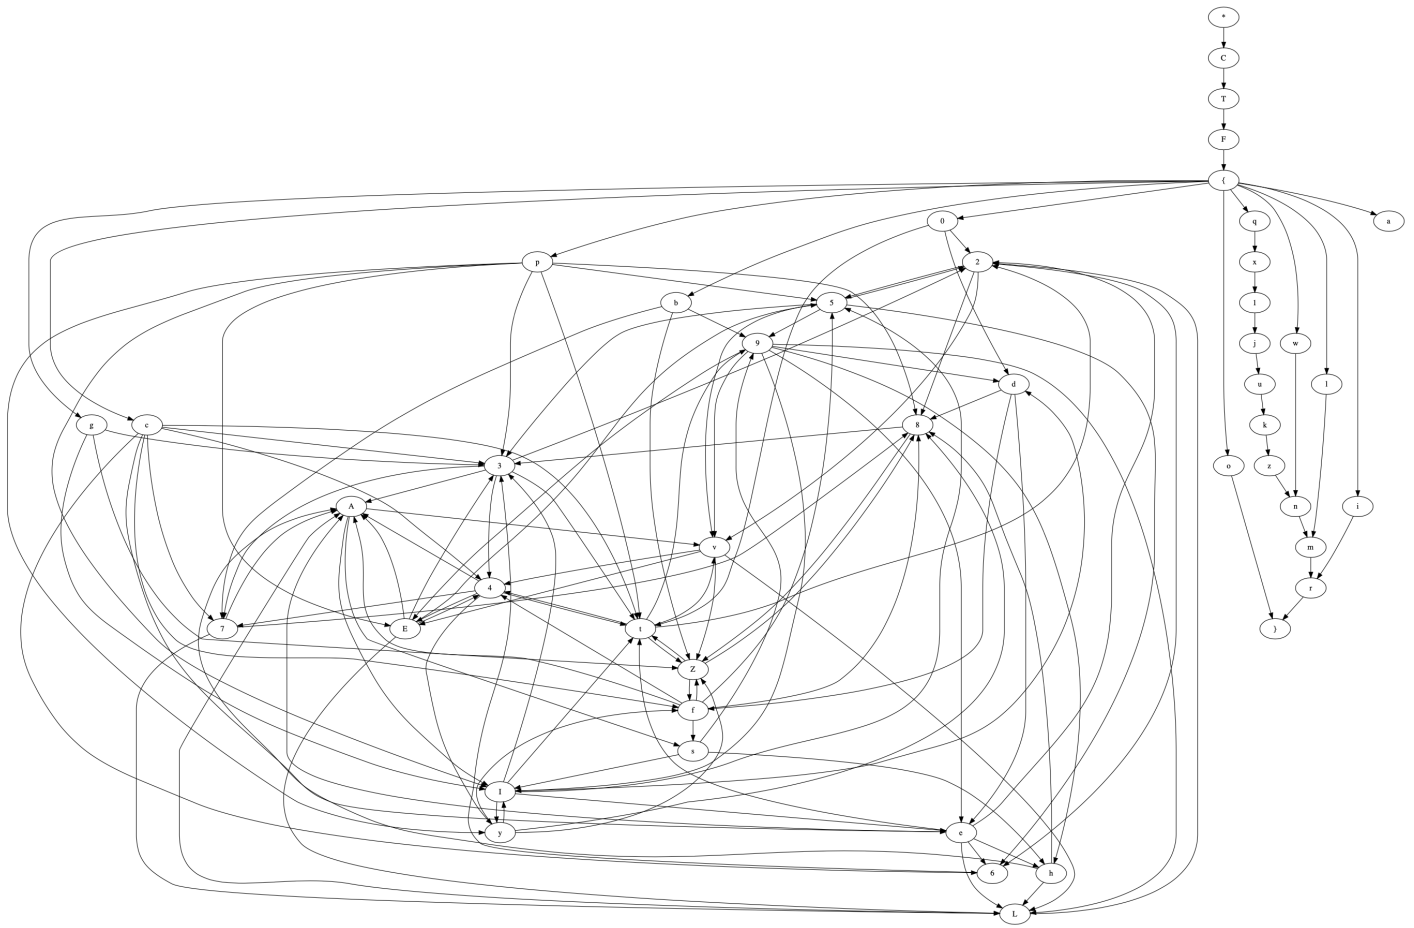
Effectively, this model maps a 47-size input vector to a 1-size output vector. Coincidentally, our available alphabet has 47 characters and therefore we can likely deduce that our input should be one-hot encoded.

However, given an input matrix of size $[47, 16]$ like we might expect, each individual character will be processed separately and therefore, this network couldn't implement flag checking.

We notice that one of the weight matrices, v_0 , contains only 0's and 1's despite being a float vector:



Furthermore, we notice an interesting pattern near the part of the alphabet with `*CTF` that indicates this might represent the transition function of a DFA. Upon plotting the transitions, we obtain:



On the right hand side, we see a clear path from `*CTF` to `}` that matches our expected size of 16: `*CTF{qx1jukznmr}`

misc/Alice's challenge

In this challenge we are provided with `Net.model`, a serialized pytorch model and 25 gradient tensors in `grad.rar`.

Looking at the model, we see three convolutional layers followed by a linear layer. Although it is not present in the serialized model, we can deduce that there is a flattening step between the two and therefore we can reconstruct the model source:

```

1 class AliceNet2(torch.nn.Module):
2     def __init__(self):
3         super().__init__()
4         self.conv = Sequential(
5             Conv2d(3, 12, kernel_size=(5,5), stride=(2,2), padding=(2,2)),
6             Sigmoid(),
7             Conv2d(12, 12, kernel_size=(5,5), stride=(2,2), padding=(2,2)),
8             Sigmoid(),
9             Conv2d(12, 12, kernel_size=(5,5), stride=(1,1), padding=(2,2)),
10            Sigmoid(),
11            Conv2d(12, 12, kernel_size=(5,5), stride=(1,1), padding=(2,2)),
12            Sigmoid(),
13        )
14        self.fc = Sequential(
15            Linear(768, 200)
16        )
17
18    def forward(self, x):
19        z0 = self.conv(x)
20        z1 = torch.reshape(z0, (-1,))
21        z2 = self.fc(z1)
22        return z2

```

With gradient tensors, our objective is likely to recover the original training samples that produced such values.

I.e. given F_θ and gradient Δ_k find x_k, y_k such that $L = \|F_\theta(x) - y\|_2$ and $\frac{\partial L}{\partial \theta} = \Delta_k$. In order to recover x_k and y_k , we can create a new loss function: $L' = \|\frac{\partial L}{\partial \theta} - \Delta_k\|_2$ and perform standard gradient descent on the input pair (x_k, y_k) .

This concept is implemented in the following pytorch code:

```

1  model = torch.load('./Net.model', map_location=torch.device('cpu'))
2
3  tensors = [
4      torch.load('./grad/%d.tensor' % x, map_location=torch.device('cpu'))
5      for x in range(25)
6  ]
7
8  def mse(a,b):
9      return torch.sum((a-b)**2)
10
11 def compute_grad_loss(x0,y0,t):
12     y1 = model(x0)
13
14     loss = torch.mean((y1 - y0) ** 2)
15     loss.backward(retain_graph=True)
16
17     grad_loss = 0
18     for a,b in zip(t, model.parameters()):
19         b0 = torch.autograd.grad(loss, b, retain_graph=True, create_graph=True)[0]
20         grad_loss += mse(a,b0)
21
22     # Compute gradient w.r.t our input pair (x,y)
23     dx = torch.autograd.grad(grad_loss, x0, retain_graph=True)[0]
24     dy = torch.autograd.grad(grad_loss, y0, retain_graph=True)[0]
25
26     return dx, dy, grad_loss
27
28 def solve_image(t):
29     x = torch.zeros(1,3,32,32)
30     y = torch.zeros(1,200)
31
32     alpha = 0.1
33
34     for i in tqdm(range(500)):

```

```

35     x0 = x.clone()
36     y0 = y.clone()
37     x0.requires_grad = True
38     y0.requires_grad = True
39
40     dx, dy, grad_loss = compute_grad_loss(x0, y0, t)
41
42     x -= dx * alpha
43     y -= dy * alpha
44
45     return x
46
47 x = [solve_image(t) for t in tensors]

```

We obtain pretty clear original images after only 500 steps:



*CTF{PZEXHBEARAK8MQT5NAliceE}

misc/babyFL